

Helical Slug Streams in an Inflated Fabric Torus: A Control-Capable Architecture for Active-Support Orbital Rings

Abstract

This paper proposes an orbital-ring architecture built from magnetically guided high-speed slug streams running in small-helix-angle lanes around a very large lightweight fabric torus. The paper's main claim is not that such a ring is presently practical. It is that this architecture introduces two linked ideas that are genuinely useful at the concept level.

The first is that helical slug streams can do more than circulate momentum. By forcing the streams to follow curved helical paths on a large toroidal membrane, the design converts momentum redirection into distributed outward pressure, hoop prestress, and local structural rigidity. The result is not a passive rigid pipe. It is a prestressed active membrane structure, more like a giant fabric sock or fabric torus wrapped around the world, whose local stiffness is created by internal moving mass.

The second is that the same helical geometry enables a four-lane balanced cell that can generate distributed tug fields for macro-scale control of the orbital ring. Spatial acceleration and deceleration of selected lanes can create bending moments and long-wavelength alignment authority without requiring literal fill-and-drain stations at ordinary control points.

A central motivation for this architecture is that simple straight-lane or monolithic-rotor orbital-ring concepts are not passively stable

under curvature perturbations. A fast mass stream pushes into existing curvature. Without a surrounding structure and active control system that can react against that tendency, perturbations are anti-restored rather than damped away. The helical fabric-torus architecture is therefore not cosmetic. It is a structural answer to the instability problem.

After presenting the force mechanism, lane-level guide requirements, perturbation argument, helical membrane architecture, and four-lane macro-control scheme, the paper turns to feasibility screens. Those screens remain severe. Full-scale macro lift requires superorbital ring-tangential momentum flux, failures are energetically extreme, and losses, thermal rejection, and containment all feed back into passive structural weight. The architecture should therefore be understood as a new way to make an orbital ring controllable and structurally legible in principle, not as evidence that an orbital ring is easy.

1. Introduction

The usual picture of an orbital ring is seductively simple: put a very fast moving mass around Earth, let curvature redirect momentum, and use the resulting reaction force to support a ring. But that picture leaves out two problems that are not secondary.

First, a fast moving mass stream is not passively self-stabilizing. If its path develops a curvature perturbation, the stream pushes further into that curvature. A simple straight lane or monolithic rotor therefore does not merely need support force. It needs a surrounding architecture that can resist and control a fundamentally anti-restoring tendency.

Second, even if one has enough moving momentum to loft a ring, that does not by itself provide a good mechanism for local structural rigidity or macro-scale alignment control. A ring around Earth has to survive construction tolerances, tether loads, payload impulses, gravity-gradient effects, and long-wavelength wobble. A concept that only says "there is a rotor going around the planet" has not yet explained how the machine is to be controlled.

This paper proposes a specific answer to both problems.

1. **Helical small- α slug streams in a lightweight fabric torus.** The lanes are wrapped helically around a very large tensile membrane tube. Their curvature creates outward pressure that inflates and prestresses the torus, turning moving momentum into local structural stiffness.
2. **A four-lane balanced cell for macro-scale control.** The same helical geometry allows balanced groups of lanes whose momentum components cancel in steady operation but can be modulated to produce distributed tug fields and ring-scale bending moments.

Those are the two main novelties of the paper.

Suggested Figure 1. Overall concept sketch. Show a world-encircling torus with roughly 100 m diameter, rendered as a lightweight fabric membrane rather than a solid pipe. Include a local cutaway showing several shallow-angle helical lanes attached to the membrane and one enlarged inset showing a four-lane balanced cell. The image should make the paper's core claim visible in one glance: the same helical lane system gives both local prestress and macro-control authority.

The argument proceeds in six steps. First, momentum redirection is established as the basic force mechanism. Second, the paper asks what one lane demands from its magnetic guide. Third, it shows why a straight high-speed lane is not passively stable under perturbation. Fourth, it introduces the helical fabric-torus architecture as a way to turn that problem into useful local prestress. Fifth, it develops the four-lane balanced cell and distributed tug fields as the main macro-control result. Only then does it turn to the harder practicality screens: macro lift throughput, containment, losses, and thermal closure.

2. Momentum redirection is the force mechanism

Consider a mass stream moving at speed v along a guide path of curvature κ . The stream requires normal acceleration

$$a_n = v^2 \kappa.$$

For a continuous stream of line density μ , the required guide force per unit path length is

$$f = \mu v^2 \kappa.$$

For a discrete slug train with fixed mass flux $\dot{m}$, the effective line density is $\mu_{\text{eff}} = \dot{m}/v$, so the corresponding force law becomes

$$f = \dot{m} v \kappa.$$

That distinction matters. A continuously filled cable-like stream at fixed line density scales as v^2 . A throughput-limited slug train whose spacing changes with speed scales as v .

The physical interpretation is straightforward. The guide pushes the stream onto a curved path. The stream pushes back on the guide

with equal and opposite force. This is the structural force source. No speculative physics is required.

It is often convenient to introduce an equivalent dynamic-tension scale

$$T_{\text{eq}} = \mu v^2$$

for the continuous case, or equivalently $T_{\text{eq}} = \dot{m}v$ for the fixed-flux slug-train case. That notation is useful, but it must not be over-read. In this architecture the support force does not require a passive material hoop carrying literal mechanical tension equal to T_{eq} . The force can instead arise from guided momentum redirection in discrete moving masses.

Suggested Figure 2. Momentum-redirection force law. Draw one curved guide segment with slugs moving along it. Show the tangent velocity, the local radius of curvature, the inward guide force on the slug, and the equal outward reaction on the guide. Include both labels $f = \mu v^2 \kappa$ and $f = \dot{m}v \kappa$ so the reader sees immediately that the same geometry applies to continuous and discrete pictures.

3. What one lane requires from its guide

Before discussing orbital-ring architecture, it is worth asking what one lane demands from its stator.

3.1 Magnetic reaction scale

A rough upper bound on magnetic normal stress is

$$p_{\text{mag,max}} \sim \frac{B^2}{2\mu_0},$$

where B is field strength and μ_0 is the permeability of free space. Ideal field-pressure scales are therefore on the order of 0.10 MPa at 0.5 T, 0.40 MPa at 1 T, 1.59 MPa at 2 T, and 3.58 MPa at 3 T. Real delivered traction will be lower because of gap, fringing, force margin, thermal limits, imperfect field topology, and control requirements.

If A' is effective magnetic interaction area per unit lane length, then the required mean traction is

$$p_{\text{req}} = \frac{f}{A'}.$$

For a fixed-flux slug lane this becomes

$$p_{\text{req}} = \frac{\dot{m}v\kappa}{A'}.$$

So tight curvature, high speed, and small interaction perimeter all make the lane harder to guide.

3.2 Convective local control

A lane is not merely static curvature plus average force. The stream is convected through the guide at high speed, so disturbances arrive on a timescale

$$t_{\text{conv}} \sim \frac{\lambda}{v},$$

where λ is disturbance wavelength. At $v = 10$, km/s, a 100 m disturbance convects past in 10 ms, a 10 m disturbance in 1 ms, and a 1 m disturbance in 0.1 ms. High speed helps force generation and hurts control.

Using lateral displacement $y(x, t)$ along a lane coordinate x , the convective derivative is

$$\frac{D}{Dt} = \frac{\partial}{\partial t} + v \frac{\partial}{\partial x},$$

so the lateral acceleration contains

$$\frac{D^2 y}{Dt^2} = y_{tt} + 2vy_{xt} + v^2 y_{xx}.$$

The last term is the same curvature-following term that generates support force. The middle term is transport coupling. Both must be handled by the guide.

A credible architecture therefore needs layered control:

- very fast local gap control at the lane level,
- bus-level coordination of nearby lanes and sectors,
- and slower macro-shape control at the ring scale.

The orbital ring cannot be treated as a single centralized control loop.

Suggested Figure 3. One lane as a guided control problem.

Show a single lane with magnetic stator modules, guide gap, slug train, sensors, and fast local controllers. Include a disturbance entering from the left and being convected downstream. The figure should make it obvious that the lane problem is distributed and time-critical, not a quasi-static guideway.

4. Why straight lanes and monolithic rotors are not passively stable

This is the missing structural argument in many simple orbital-ring pictures.

Consider a nominally straight lane or monolithic rotor segment with a small transverse deflection $y(x, t)$. For small slope, the centerline curvature is approximately

$$\kappa \approx y_{xx}.$$

The moving stream must then be forced to follow that curved path. The stream exerts on the guide a reaction load per unit length of the form

$$q_{\text{stream}} \approx -T_{\text{eq}} y_{xx},$$

where $T_{\text{eq}} = \mu v^2$ or $T_{\text{eq}} = \dot{m}v$, depending on whether one uses the continuous or fixed-flux picture.

Now consider a sinusoidal perturbation,

$$y(x, t) = Y \sin(kx).$$

Then

$$y_{xx} = -k^2 y,$$

so the stream reaction becomes

$$q_{\text{stream}} = +T_{\text{eq}} k^2 y.$$

That has the **same sign as the displacement**. If the lane is displaced upward at some point, the stream pushes upward there as well. The moving mass therefore pushes further into an existing curvature perturbation rather than restoring the lane to straightness.

This is not yet a complete stability model. Real dynamics also involve beam stiffness, membrane stiffness, damping, actuator delay, and control law. But it is enough to establish the key point: a straight high-speed lane is not passively self-centering. The moving stream behaves like an anti-restoring curvature follower unless the surrounding structure and controls provide compensating reaction.

That observation cuts directly against naive pictures of a single giant monolithic rotor or a set of parallel straight rotors simply going around Earth. Such concepts may contain enough momentum to generate support in principle, but they do not automatically contain a good structural answer to perturbations.

Suggested Figure 4. Anti-restoring behavior of a perturbed straight lane. Show a nominally straight lane with a small sinusoidal bend. Draw the local curvature and the stream reaction arrows pushing into the bend rather than away from it. Under the drawing, show the short derivation $q_{\text{stream}} \approx -T_{\text{eq}}y_{xx} \rightarrow +T_{\text{eq}}k^2y$. This figure should be the moment where the reader understands why a simple straight rotor picture is incomplete.

5. Distributed control and why fill-and-drain is not required

If macro-control is to come from lane-speed modulation, the next question is whether changing speed requires literal insertion and removal of slugs at every control station. For ordinary control sections, the answer appears to be no.

Let slug headway be h . In a locally uniform region with lane speed v , the center-to-center spacing is

$$s = vh.$$

A commanded speed field changes occupancy by compressing or expanding that spacing. That means useful control can be produced by distributed acceleration and deceleration zones without ordinary fill-and-drain hardware.

For monotone slowdown from v_0 to $v_1 < v_0$, the minimum spacing is

$$s_{\min} = hv_1,$$

so collision avoidance requires

$$hv_1 \geq \ell_s + g_{\min},$$

where ℓ_s is slug length and $g_{\min}$ is minimum allowed gap.

For an upward ramp, incomplete actuation can let a trailing slug catch a slower incumbent. A useful first delay bound is

$$\tau_d < \frac{hv_0 - (\ell_s + g_{\min})}{v_1 - v_0},$$

where τ_d is total sensing, computation, actuation, and field-establishment delay.

The result is conceptually important. A lane can support smooth distributed transition zones in which slug spacing changes kinematically. That means the macro-control architecture can be based on continuous tug fields rather than on exotic station hardware every time one wants to make a moment.

Suggested Figure 5. Distributed speed transition without fill and drain. Show a slug train entering a gradual slowdown region, with spacing shrinking smoothly through the transition and then becoming uniform again. Add a second panel with a speed-up region. The point

is to let the reader see that occupancy change can occur through spacing evolution alone.

6. Helical lanes in a lightweight fabric torus

The helical fabric-torus architecture is the first major payoff of the paper.

Instead of asking a fast lane to remain straight in free space, the architecture places the lane on an intentionally curved helical path wrapped around a very large tensile membrane tube. The tube should be imagined as a lightweight fabric torus, likely made from a high-tensile membrane family rather than a heavy rigid shell. "Fabric sock around the world" is crude language, but it is not the wrong mental image.

6.1 Helical curvature turns instability into pressure

For a helix on a cylinder of radius a , with helix angle α measured relative to the tube axis,

$$\kappa_{\text{helix}} = \frac{\sin^2 \alpha}{a}.$$

That curvature is not an accident. It is the local mechanism that converts moving momentum into distributed outward reaction on the membrane-supported lane system.

For a fixed-flux slug lane, the local outward load per unit axial ring length is

$$q_{\text{loc, lane}} = \dot{m}u \frac{\tan^2 \alpha}{a},$$

where

$$u = v \cos \alpha$$

is the ring-tangential or tube-axial component of lane speed.

For one symmetric pair of lanes,

$$q_{\text{loc,pair}} = 2\dot{m}u \frac{\tan^2 \alpha}{a}.$$

Spread over cylindrical area $2\pi a$ per unit axial length, one pair produces equivalent pressure

$$p_{\text{pair}} = \frac{\dot{m}u \tan^2 \alpha}{\pi a^2}.$$

For N_p paired modules distributed around the torus,

$$p_{\text{eq}} = \frac{N_p \dot{m}u \tan^2 \alpha}{\pi a^2}.$$

So the helical lanes actively inflate the torus.

6.2 Membrane hoop tension creates a locally stiff substrate

Equivalent pressure p_{eq} creates hoop membrane force per unit axial length

$$N_\theta = p_{\text{eq}} a.$$

This prestress plausibly helps with:

- maintaining circular cross-section,
- resisting ovalization,
- suppressing wrinkling in a membrane wall,
- increasing local indentation stiffness,

- and providing a stable substrate onto which guides, sensors, power hardware, and auxiliary structure can be mounted.

That is the architectural move. The lane curvature that would be destabilizing in a nominally straight free-space guide is now deliberately redirected into useful membrane prestress.

Put differently, the lane is no longer trying to define its own path in empty space. It is mounted to a prestressed toroidal membrane whose tensile strength provides the reaction structure. Qualitatively, the membrane-plus-guide assembly acts like a structural potential well around the intended helical path.

Lanes could in principle be mounted on the inner or outer surface of the membrane, or on shallow truss or rib hardware attached directly to it. The point is not the exact attachment detail at this stage. The point is that the inflated membrane becomes the thing the lane pushes against.

Suggested Figure 6. Helical lane inflating a fabric torus. Draw a cutaway of the torus wall with shallow-angle helical lanes attached to it. Show the lane reaction pressing outward, the membrane carrying hoop tension, and secondary hardware mounted to the now-prestressed wall. The figure should make clear that the membrane is not decorative cladding. It is the structural reaction surface for the helical lanes.

6.3 Why the helix angle is likely small

The architecture wants helical curvature for local prestress, but it also wants most of the speed to remain ring-tangential so the orbital ring can loft itself. That trade pushes the design toward shallow helical angles.

Define the aggregate ring-tangential momentum-flux scale

$$A = N_p \dot{m} u.$$

The inflation requirement can be written as

$$A \geq \frac{\pi a N_{\theta, \text{req}}}{\tan^2 \alpha},$$

where $N_{\theta, \text{req}}$ is required hoop membrane force per unit length.

The macro-lift requirement, derived later, is

$$A \geq \frac{w_p}{2 \left(\frac{1}{R} - \frac{g_h}{u^2} \right)},$$

where w_p is passive weight per unit ring length, R is ring radius, and g_h is gravity at altitude.

Balancing the two gives a useful screening estimate

$$\tan^2 \alpha_{\text{opt}} = 2\pi a \Gamma \left(\frac{1}{R} - \frac{g_h}{u^2} \right), \quad \Gamma = \frac{N_{\theta, \text{req}}}{w_p}.$$

For small angles,

$$\alpha_{\text{opt}} \approx \sqrt{2\pi a \Gamma \left(\frac{1}{R} - \frac{g_h}{u^2} \right)}.$$

Using a screening case with $a = 50$, m, altitude 80 km, $u = 10$, km/s, and $\Gamma = 10$, one gets

$$\alpha_{\text{opt}} \approx 0.78^\circ.$$

That is a striking result. Once macro lift matters, the optimal helical bias may be very small. This supports exactly the picture advocated

here: a large torus with many shallow-angle helical lanes, not a steeply wrapped screw conveyor.

6.4 What the inflated torus does and does not solve

The inflated fabric torus solves an important local problem. It gives the lanes a reaction structure, creates local rigidity, and provides a platform for hardware.

It does **not** automatically create a globally rigid torus. A 100 m diameter prestressed membrane can be locally round and still be very flexible at long wavelength. The orbital ring still requires active macro-control.

That second problem is where the four-lane balanced cell becomes essential.

7. The four-lane balanced cell

A simple mirrored pair is useful, but it is not enough for a helical architecture.

If one chooses a mirrored helical pair with tangent vectors

$$\begin{aligned}t_x &= \cos \alpha, e_z + \sin \alpha, e_\theta, \\t_y &= -\cos \alpha, e_z + \sin \alpha, e_\theta,\end{aligned}$$

then the axial components cancel, but the circumferential components add. That means a two-lane pair can still carry steady circumferential momentum and angular momentum around the torus.

The minimal balanced cell therefore contains four lanes:

1. right-handed helix, positive axial travel,

2. left-handed helix, negative axial travel,
3. left-handed helix, positive axial travel,
4. right-handed helix, negative axial travel.

Operated together, this cell can cancel to first order:

- net axial momentum,
- net circumferential momentum,
- net angular momentum about the torus,
- and first-order structural torque in symmetric operation.

At the same time it preserves:

- common-mode inflation pressure,
- controllable pressure trim,
- axial or tangential tug authority,
- and azimuthally selective moment generation.

This is not a decorative symmetry argument. It is the architecture that makes the helical torus usable as a control machine rather than merely a pressurized tube.

Suggested Figure 7. Four-lane balanced cell. Show the four lane helices in cross-section and in a short perspective cutaway, with arrows labeling handedness and travel direction. Use color-coded momentum vectors to show cancellation of axial, circumferential, and angular components. Then indicate the common-mode pressure channel that remains. The figure should make the cell feel like a genuine machine primitive.

8. Distributed tug fields and macro-scale ring control

This is the second major payoff of the architecture, and arguably the main result of the paper.

The inflated torus makes a locally stiff substrate. The four-lane balanced cell makes that substrate actively controllable at long wavelength.

8.1 Tug comes from speed transitions

For a mirrored pair selected to produce axial tug, a speed change Δv across a transition produces net axial tug

$$F_{\text{tug},1} = 2\dot{m}\Delta v \cos \alpha.$$

If N_s paired modules participate in one azimuthal sector,

$$F_{\text{sector}} = 2N_s\dot{m}\Delta v \cos \alpha.$$

That result is important because it comes from speed modulation alone. Ordinary control points do not need to insert or remove slugs from the lane. They need to establish distributed acceleration and deceleration fields.

If the transition is smoothed over axial width λ , the tug becomes a distributed load density

$$q_z(x) = -2N_s\dot{m} \cos \alpha, \frac{dv}{dx}.$$

Integrating across the transition gives

$$|F_z| = 2N_s\dot{m}\Delta v \cos \alpha.$$

So smoothing reduces peak local force density without destroying total tug authority.

8.2 Opposed sectors create bending moments

Apply such tug fields in opposed azimuthal sectors of a torus of radius a . Then one gets a bending moment of order

$$M \sim aF_{\text{sector}}.$$

A more explicit opposed-sector estimate with sector width $\Delta\phi$ is

$$M_{\text{pair}} = 4aN_s\dot{m}\Delta v \cos \alpha, C_{\text{sec}}(\Delta\phi),$$

with

$$C_{\text{sec}}(\Delta\phi) = \frac{\sin(\Delta\phi/2)}{\Delta\phi/2}.$$

This is the macro-control primitive. By accelerating selected lanes in one sector and compensating them elsewhere, the ring can generate controlled long-wavelength moments while keeping the cell balanced in the aggregate.

8.3 What this control channel is for

The point is not only to suppress abstract wobble. The same mechanism is the natural candidate for handling:

- long-wavelength orbital-ring alignment,
- tether loads,
- local payload-launch impulses,
- construction asymmetries,
- thermal distortion,
- and other slow or moderate disturbances that act on the ring at large scale.

A useful way to say it is this: the helical lanes pay twice. First they inflate and prestress the torus. Then, because they are grouped into balanced cells, they provide a distributed internal actuation system for macro-scale shape control.

8.4 Why this is different from a simple rotor ring

A simple monolithic rotor picture treats moving momentum mainly as a support source. The architecture proposed here treats moving momentum as both a support source and a spatially distributed control resource.

That is the novel structural-control idea the paper is trying to surface.

Suggested Figure 8. Distributed tug fields and macro-control.

Show a short ring segment with one pair of opposed azimuthal sectors highlighted. In each highlighted sector, draw gradual acceleration and deceleration zones in the participating lanes, along with the resulting axial tug vectors and net bending moment on the torus cross-section. A second panel should show the same mechanism acting on a global orbital-ring arc to resist a tether load or wobble mode.

9. What this architecture still does not solve automatically

The paper's claim is architectural novelty, not practical closure. Several severe screens remain.

9.1 Macro lift is still a separate burden

The helical torus gives local prestress. It does not, by itself, loft the entire ring around Earth. Macro lift comes from turning the ring-tangential component of stream momentum around Earth.

Let $u = v \cos \alpha$ be ring-tangential speed and let the ring radius be R . For one slug lane with mass flux $\dot{m}$, the net outward lift per unit ring length is

$$q_{\text{lift, lane}} = \dot{m} \left(\frac{u}{R} - \frac{g_h}{u} \right),$$

where g_h is gravity at altitude.

For N_p paired modules,

$$q_{\text{lift}} = 2N_p \dot{m} \left(\frac{u}{R} - \frac{g_h}{u} \right).$$

The lift changes sign at

$$u_{\text{orb}} = \sqrt{g_h R}.$$

If $u < u_{\text{orb}}$, the moving stream loads the ring downward overall. Only for $u > u_{\text{orb}}$ does it contribute net outward lift.

This is why the helical angle has to stay small at full scale. The ring needs its speed mostly in the ring-tangential direction.

9.2 The full system faces a closure loop

Let passive structural and hardware weight per unit ring length be w_p . The lift requirement can be written as

$$A = N_p \dot{m} u \geq \frac{w_p}{2 \left(\frac{1}{R} - \frac{g_h}{u^2} \right)}.$$

That sounds like a simple scaling law, but it closes several hard subsystems into one loop:

[$w_p \uparrow \rightarrow A \uparrow \rightarrow$
 $E \uparrow \rightarrow \text{kin} \uparrow \text{ and } P \uparrow \rightarrow \text{loss} \uparrow$

\rightarrow

$w_{\text{contain}}, w_{\text{thermal}}, w_{\text{power}} \uparrow$
\rightarrow $w_p \uparrow$.]

In words:

1. more passive hardware raises ring weight,
2. more weight requires more momentum flux,
3. more momentum flux raises stored energy per metre and usually raises losses,
4. those burdens demand more containment, thermal hardware, and power infrastructure,
5. which raises ring weight again.

The concept stands or falls on whether this loop converges.

9.3 Safety, losses, and thermal rejection are still existential

At one representative screening point, altitude 80 km, $u = 10, \text{ km/s}$, and $w_p = 10, \text{ kN/m}$, the required momentum-flux scale is roughly

$$A \sim 8.44 \times 10^{10}, \text{ N}.$$

That same scale is also a warning about kinetic energy per unit length. Using the fixed-flux slug-train estimate, the moving energy inventory is of the same order,

$$E'_{\text{kin}} \sim N_p \dot{m} u = A \sim 8.44 \times 10^{10}, \text{ J/m},$$

which is roughly 20 tons TNT equivalent per metre.

That is why failure localization, catcher structure, dump paths, and segment isolation are architectural requirements rather than safety afterthoughts.

Losses are similarly dangerous because radiator mass pushes back into w_p . With a rough radiative rejection capability of $q_{\text{rad}} \sim 500, \text{ W/m}^2$ and radiator areal mass around $5, \text{ kg/m}^2$, a steady loss density of 100 kW per metre implies roughly 200 m^2 of radiator area per metre and nearly $9.8, \text{ kN/m}$ of radiator weight by itself. That is enough to become a first-order structural term.

So the architecture may be conceptually strong and still prove operationally brutal.

Suggested Figure 9. Closure-loop reality check. Draw a systems loop linking passive weight w_p , required momentum flux A , stored energy per metre, steady losses, containment hardware, thermal hardware, and then back to w_p . This figure should serve as the paper's main cautionary diagram: the invention may solve the control architecture without yet solving the practicality loop.

10. A staged development path

Because the architecture combines force generation, control, prestressed membrane structure, and very high moving energy, it should be developed in stages.

1. **Bench force demonstrator.** Measure the momentum-redirection law directly.
2. **Single-lane guide demonstrator.** Show local magnetic guidance and convective control margins.
3. **Perturbation demonstrator.** Show the anti-restoring behavior of a straight lane and the stabilizing role of a surrounding reaction structure.

4. **Helical membrane demonstrator.** Show measurable inflation pressure and local stiffening in a small fabric tube.
5. **Four-lane balanced-cell demonstrator.** Show momentum cancellation with preserved pressure and control channels.
6. **Distributed-transition demonstrator.** Show that smooth speed modulation produces real tug fields without routine fill-and-drain hardware.
7. **Architectural active-tube demonstrator.** Show a visible lightweight structure whose stiffness clearly comes from internal momentum rather than passive bulk.

A public demonstrator should prioritize safety and legibility over maximum size. The right public message is not "we already built an orbital ring." It is "active internal momentum can make a lightweight membrane structure carry load and be controllable in a way passive structure alone cannot."

Suggested Figure 10. Demonstrator ladder. Draw a progression from bench guide-force rig to helical fabric tube, four-lane balanced cell, distributed-transition rig, and finally an architectural active-tube demonstrator. Label the principle proven at each step and the dangerous scale jump being intentionally deferred.

11. Conclusion

This paper argues for a specific architectural idea, not for immediate engineering closure.

The first key idea is that shallow-angle helical slug streams in a lightweight fabric torus can convert momentum redirection into distributed inflation pressure and hoop prestress. That provides a

way to make a huge membrane tube locally stiff enough to serve as the structural substrate for its own guide hardware.

The second key idea is that the same helical geometry enables a four-lane balanced cell whose lanes can be modulated to produce distributed tug fields for macro-scale orbital-ring control. That gives the ring an internal actuation system for alignment, wobble suppression, tether-load management, and useful work, while preserving first-order momentum balance.

Those two ideas belong together. The inflated torus without the balanced control cell is only a locally stiff membrane. The balanced control cell without the toroidal prestressed substrate lacks a good structural medium through which to act. Together they form a coherent architecture.

A major motivation for this architecture is that straight high-speed lanes and simple monolithic rotors are not passively self-stabilizing under curvature perturbation. The stream pushes into existing curvature. The helical fabric torus is therefore not merely a convenient geometry. It is a way to supply the reaction structure and control channels that a viable ring would need anyway.

That does not make the full orbital ring easy. Macro lift still demands superorbital ring-tangential momentum flux. Stored energy per metre is extreme. Safety, containment, losses, and thermal rejection all remain potentially decisive. The honest conclusion is therefore two-sided.

- As a **control-capable architecture**, this concept is substantially stronger than a simple rotor-around-Earth picture.
- As a **practical megastructure**, it still faces severe closure problems.

That is a useful result. Even if the full orbital ring proves too hard, the paper identifies a new family of active-support structures in which moving momentum, tensile membranes, and distributed control are tightly integrated. If the concept advances, it should advance by proving those ingredients experimentally in that order.